Deep Learning
CS60010

Abir Das
Assistant Professor
Computer Science and Engineering Department
Indian Institute of Technology Kharagpur

http://cse.iitkgp.ac.in/~adas/

Agenda

The Building Blocks of Convolutional Neural Networks/CNNs/ConvNets

Importance of MNIST

MNIST

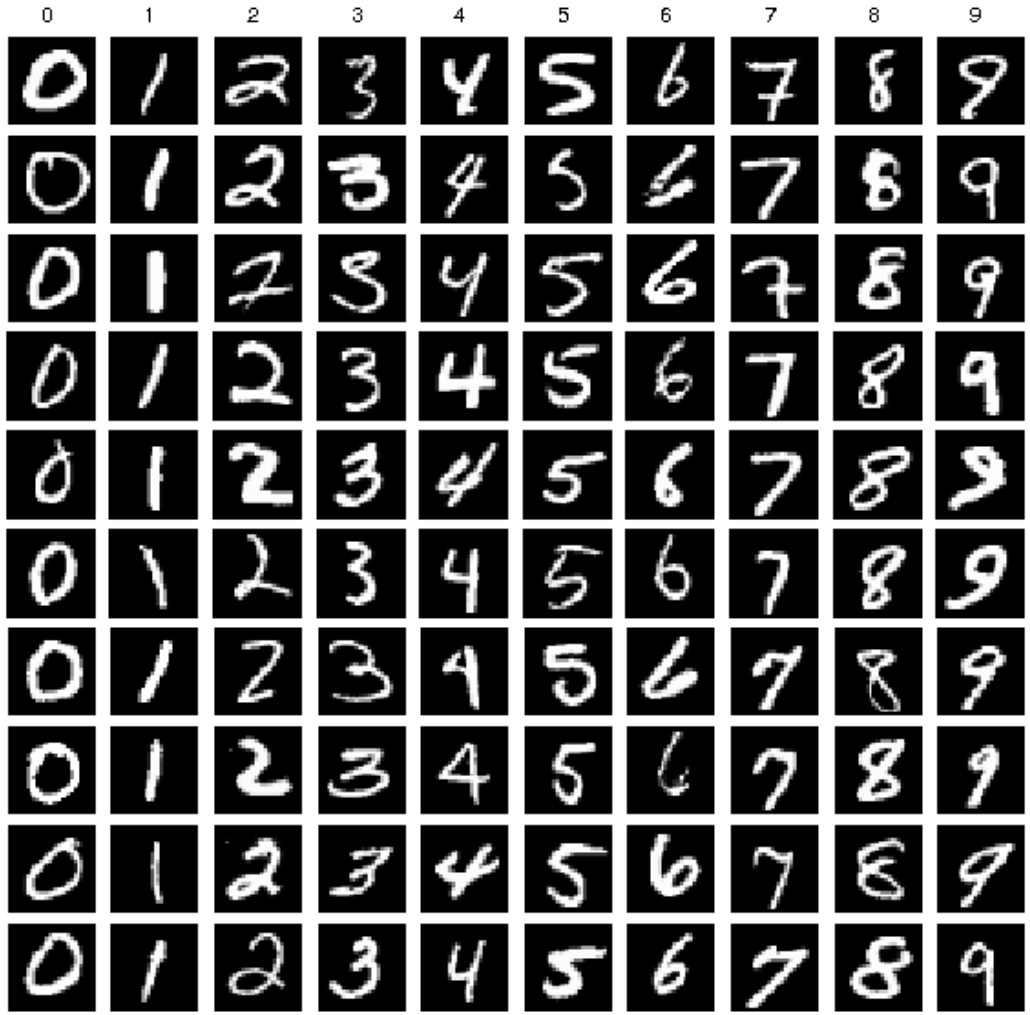

- database of handwritten digits
- 10 classes
  - Training set 60,000 images
  - Test set of 10,000 images
- Greyscale images of size 28x28
- Often treated as `Hello World' for any ML/DL practitioner

Image taken from Researchgate

Example from MNIST

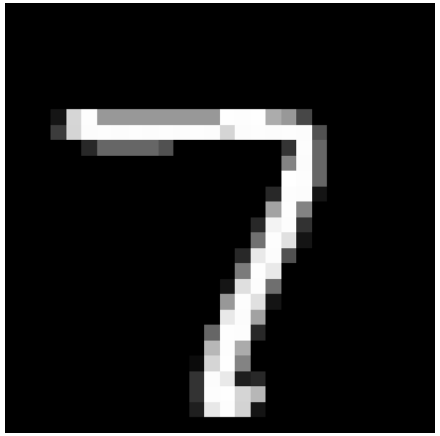

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	21	214	253	152	152	152	152	152	152	152	152	254	253	254	172	152	71	0	0	0	0	0	0	0	0
0	0	0	61	213	252	253	252	253	252	253	252	253	252	213	252	253	252	253	252	82	0	0	0	0	0	0	0
0	0	0	0	0	41	102	102	102	102	82	0	0	0	0	0	0	0	52	253	102	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	132	252	102	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	255	253	102	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	41	253	252	20	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	163	254	131	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	41	243	253	50	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	113	253	224	20	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	41	233	252	81	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	123	254	233	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	21	223	253	111	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	152	253	224	20	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	233	252	162	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	102	254	253	41	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	183	253	171	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	11	213	254	131	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	51	252	233	30	41	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	51	253	254	213	183	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	31	232	253	212	20	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

7 Like Window?

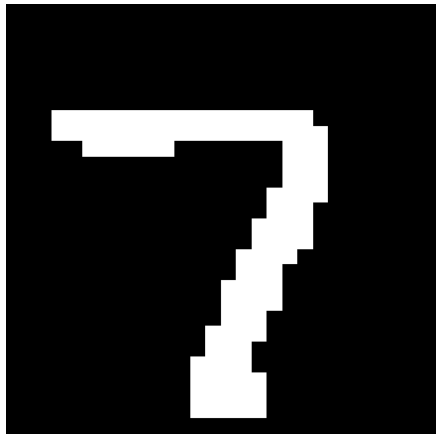

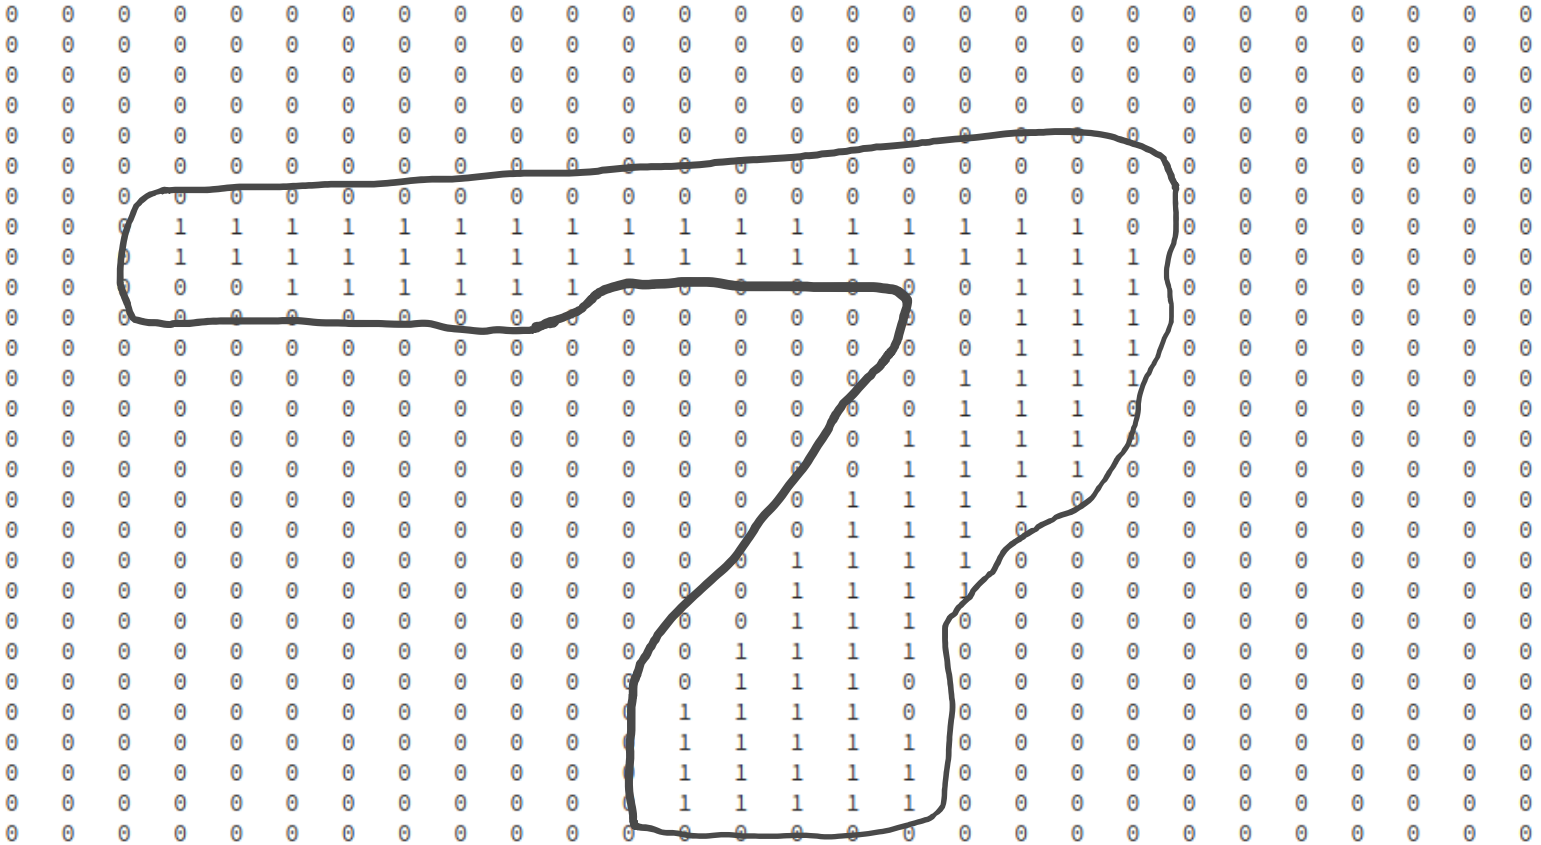

Filtering

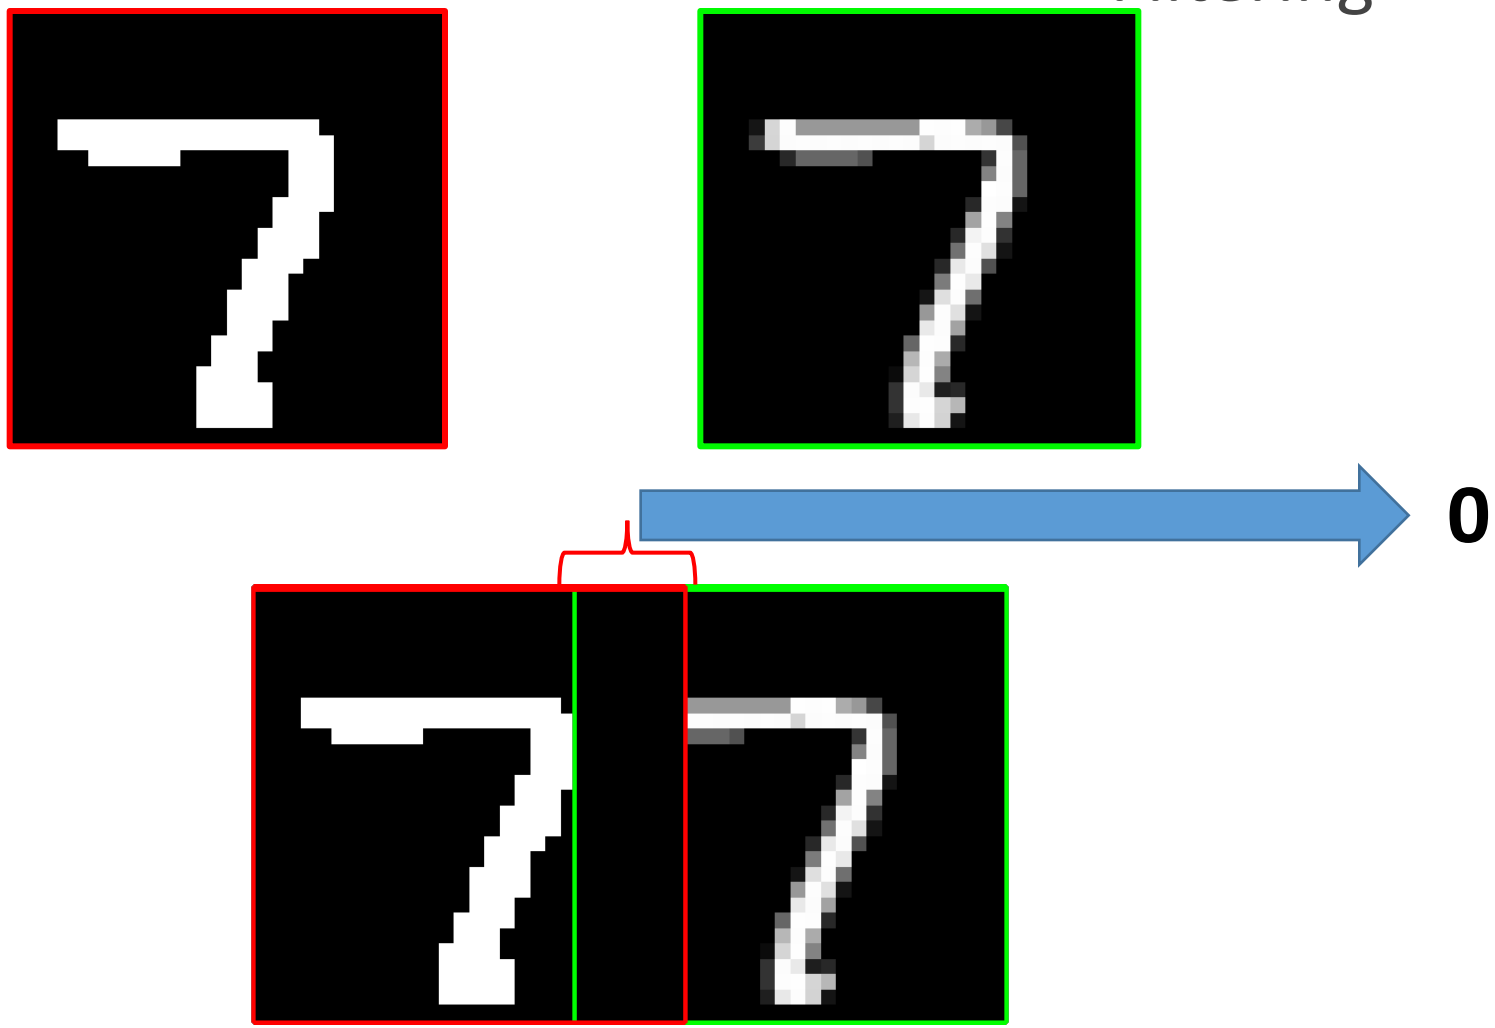

Filtering

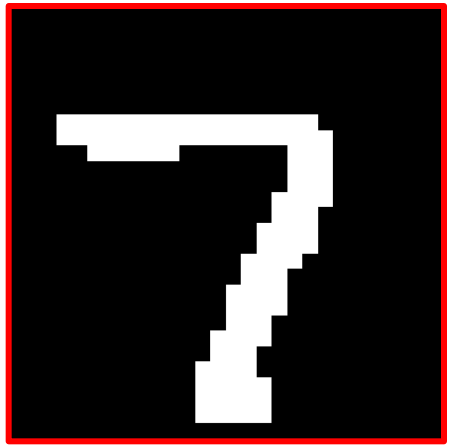

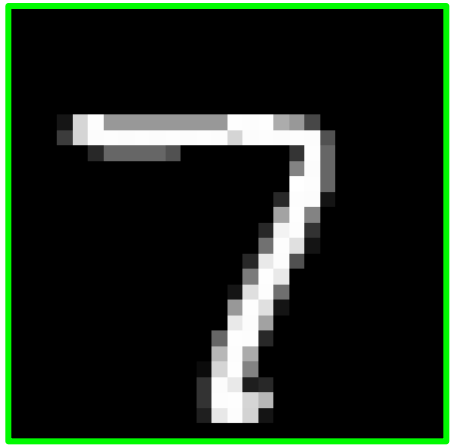

1410

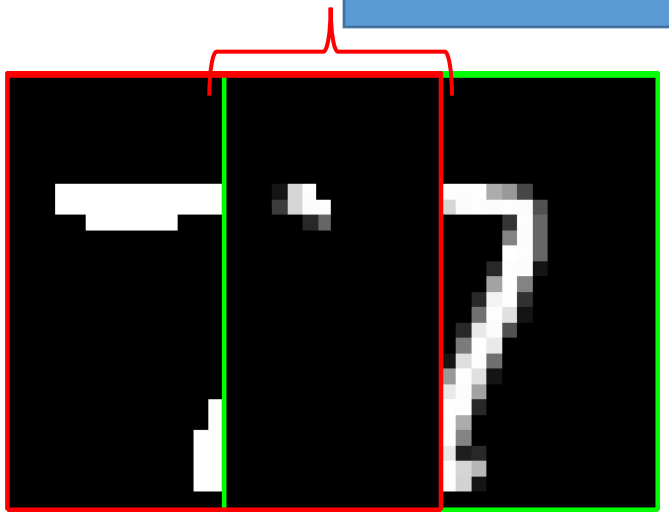

Filtering

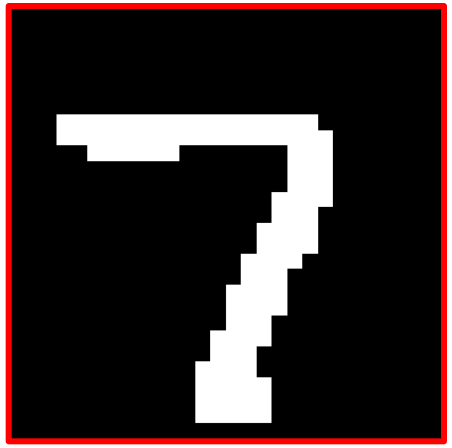

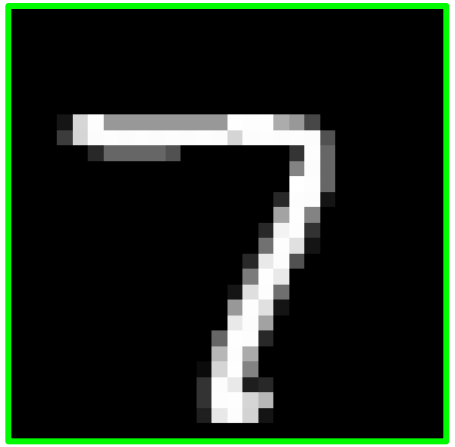

4098

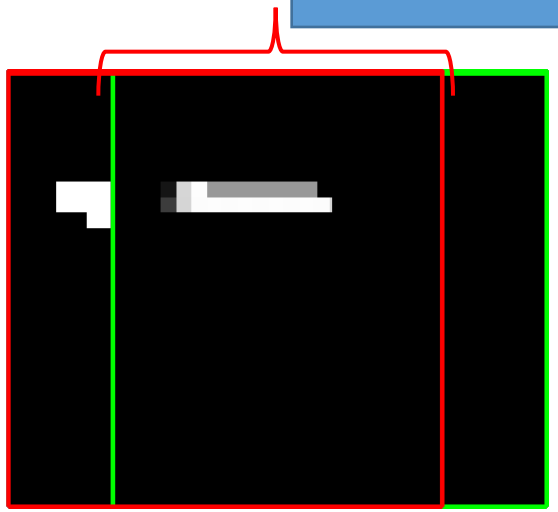

Filtering

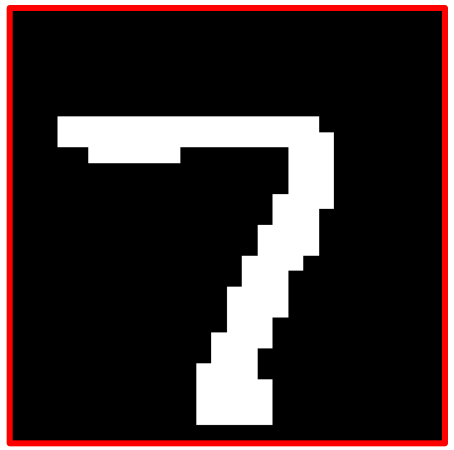

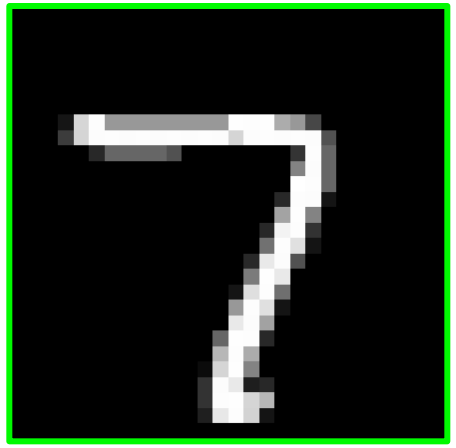

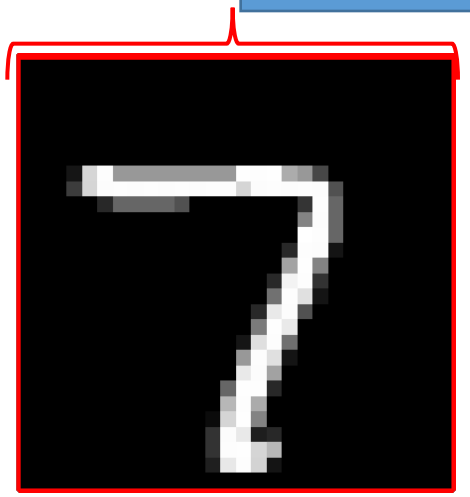

18604

Filtering

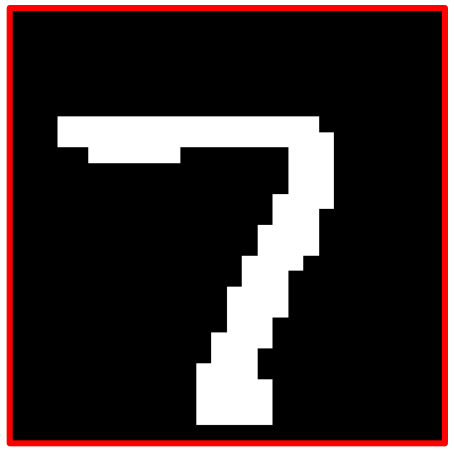

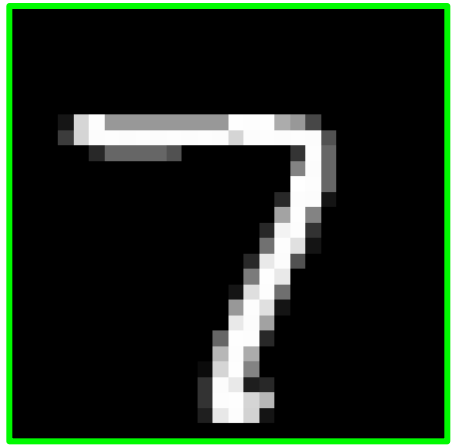

4331

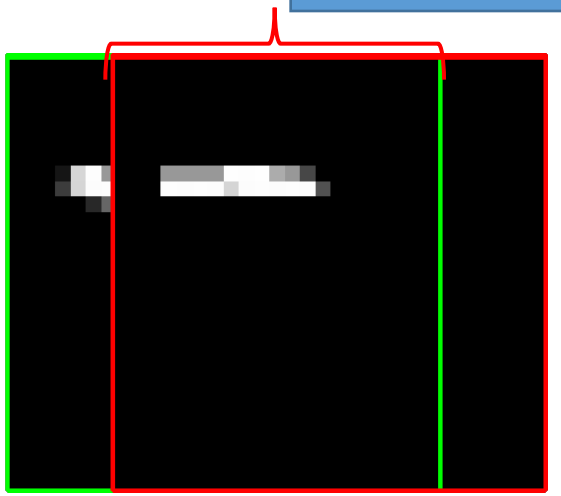

Filtering

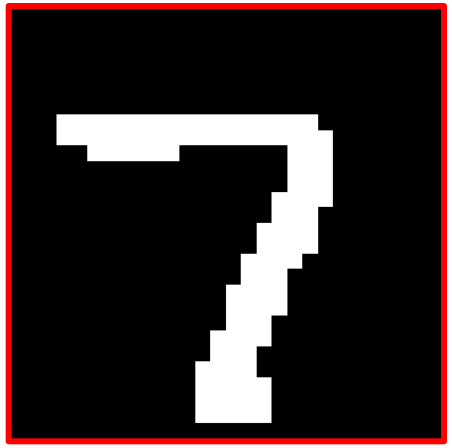

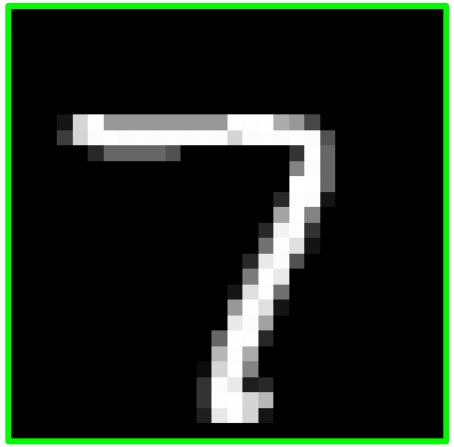

1589

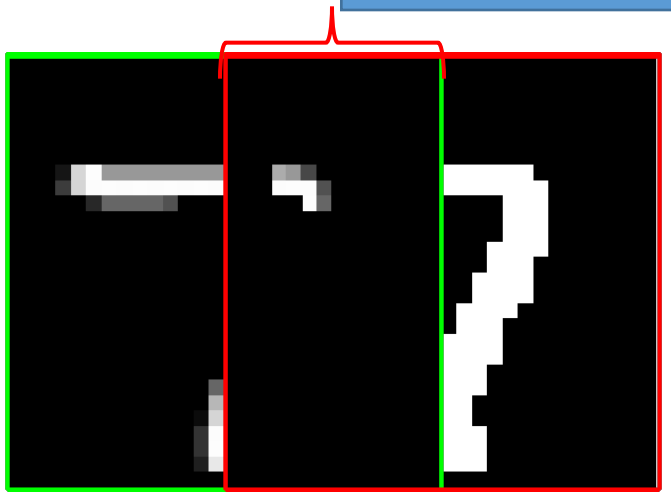

Filtering

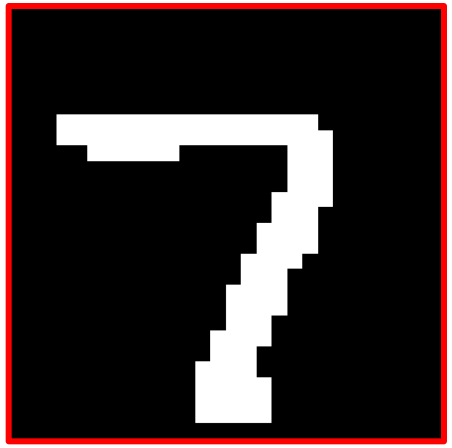

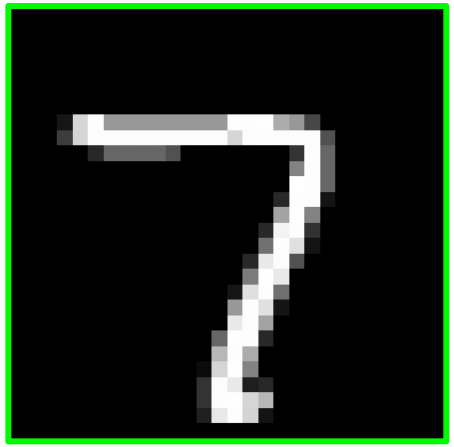

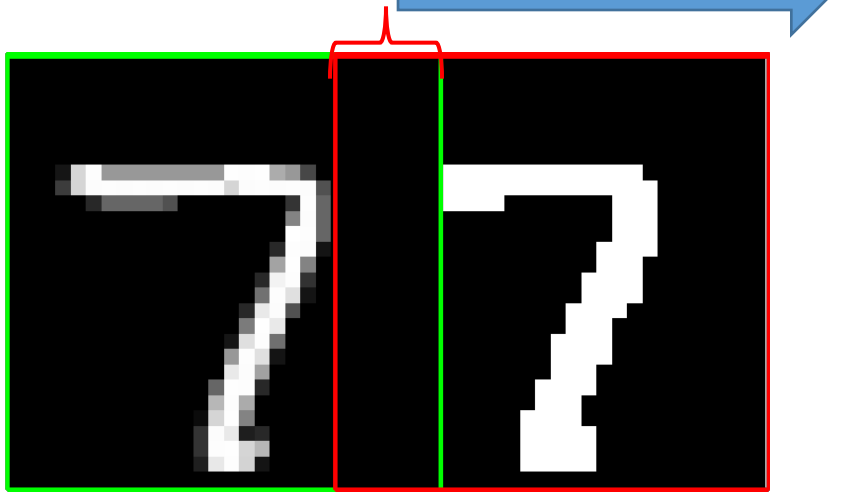

0

Classification by Matching Filters

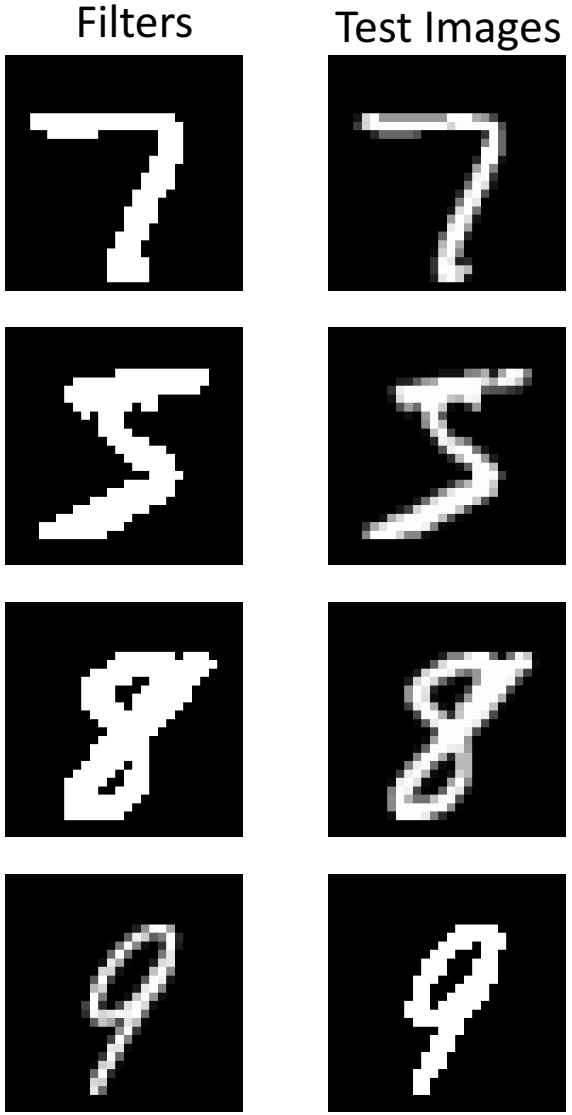

But what if the test image is a little
Rotated (or)
Skewed (or)
Zoomed out and so on

Effect of Slight Rotation

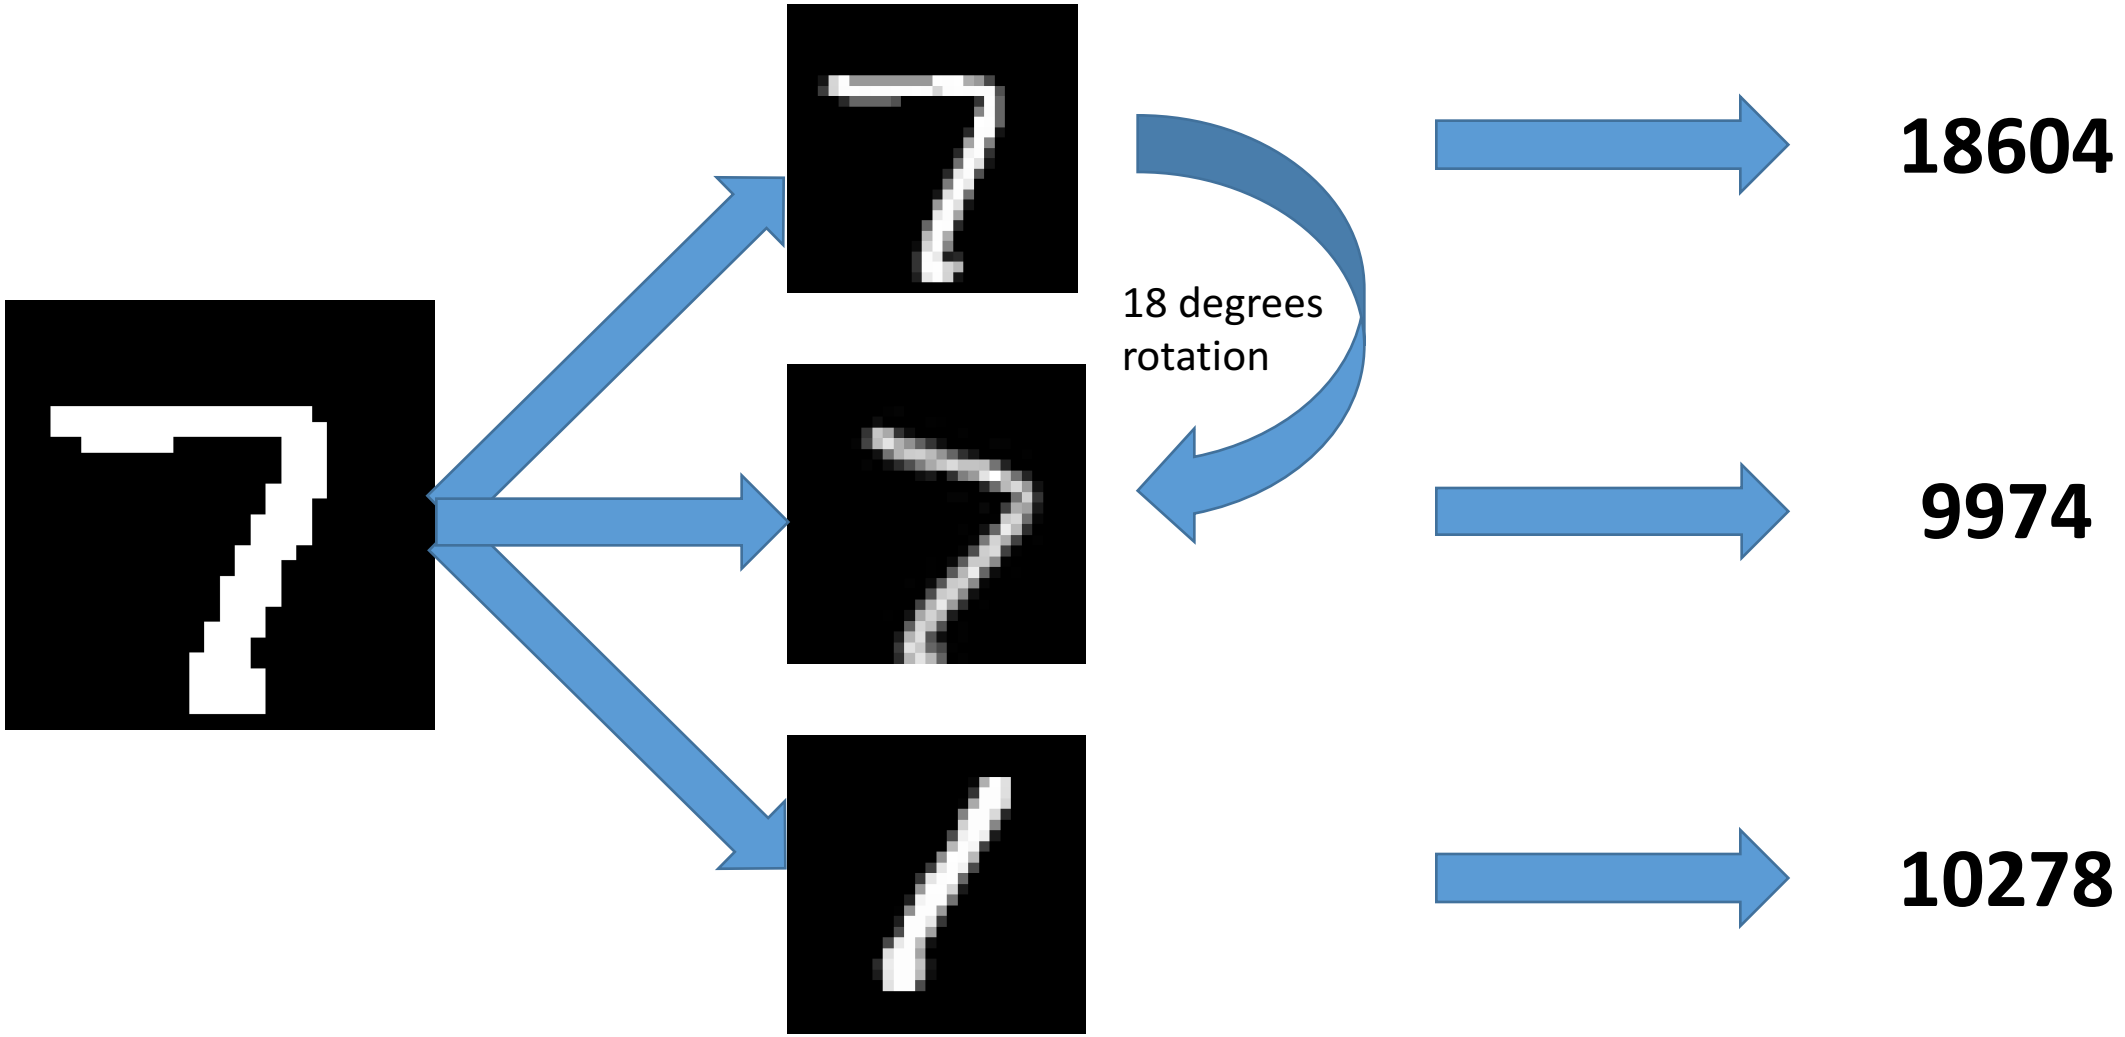

Convnets (Fukushima, LeCun, Hinton)

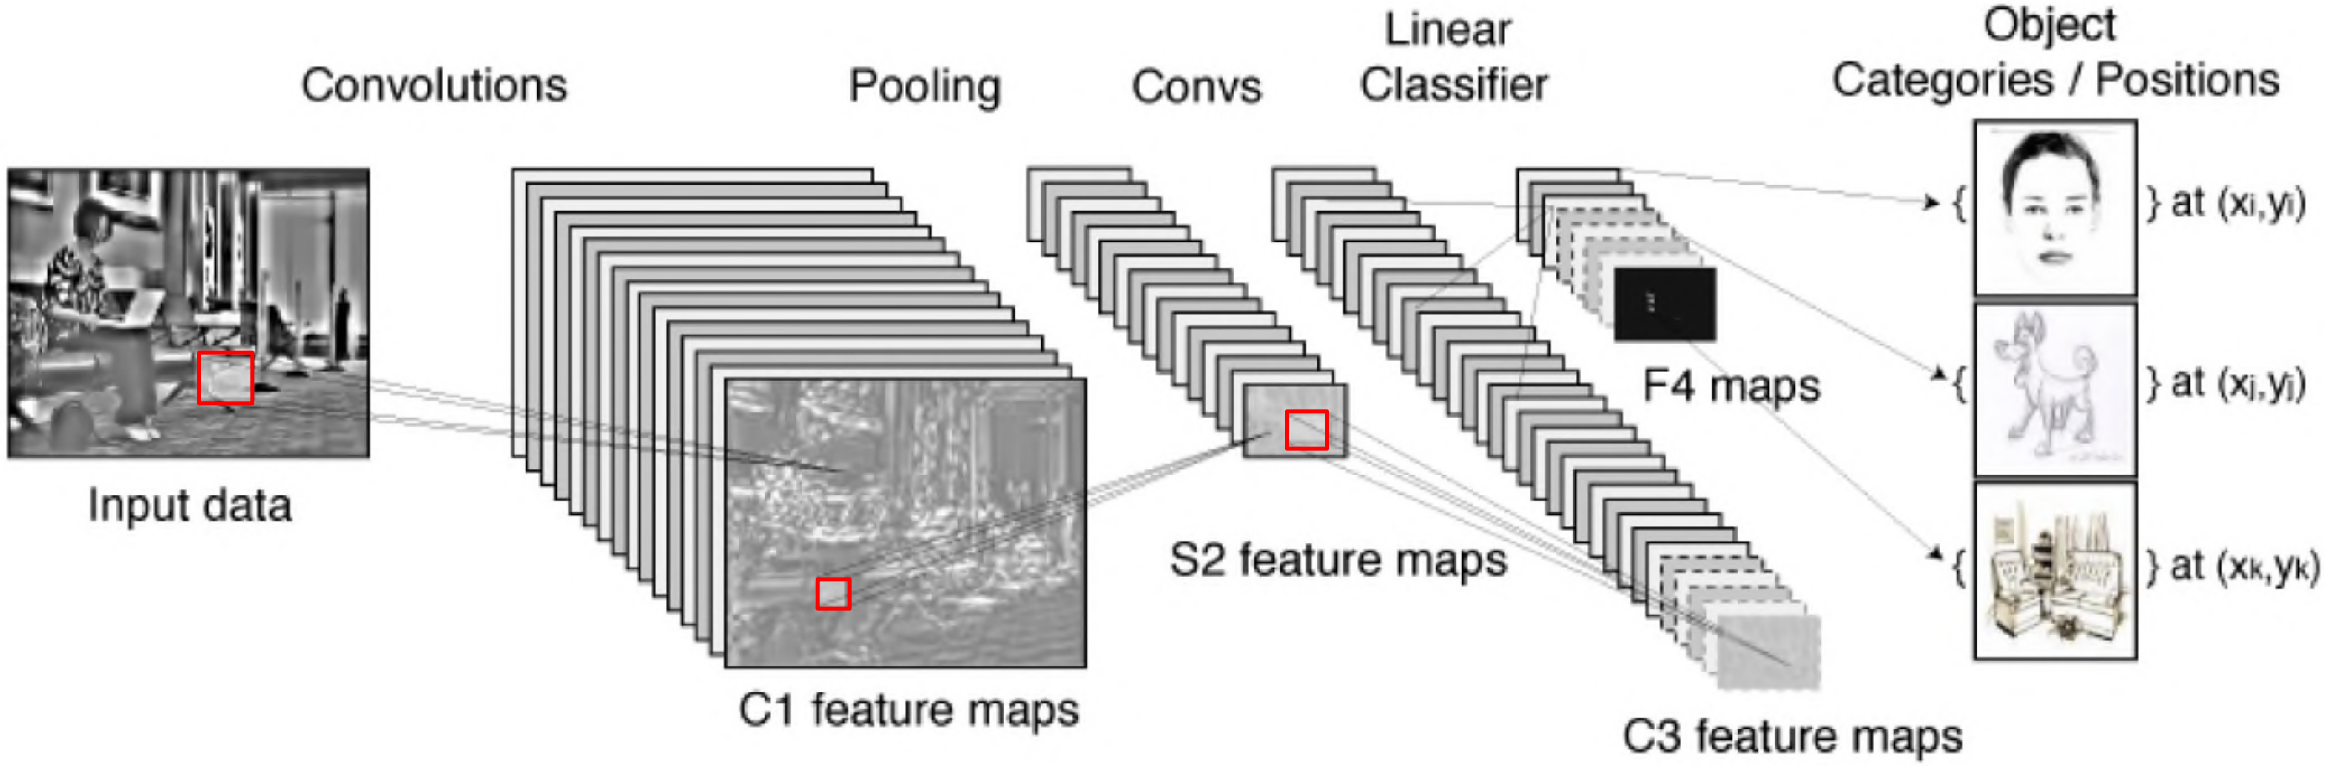

Slide courtesy: Nando de Freitas, University of Oxford

Convnets (Fukushima, LeCun, Hinton)

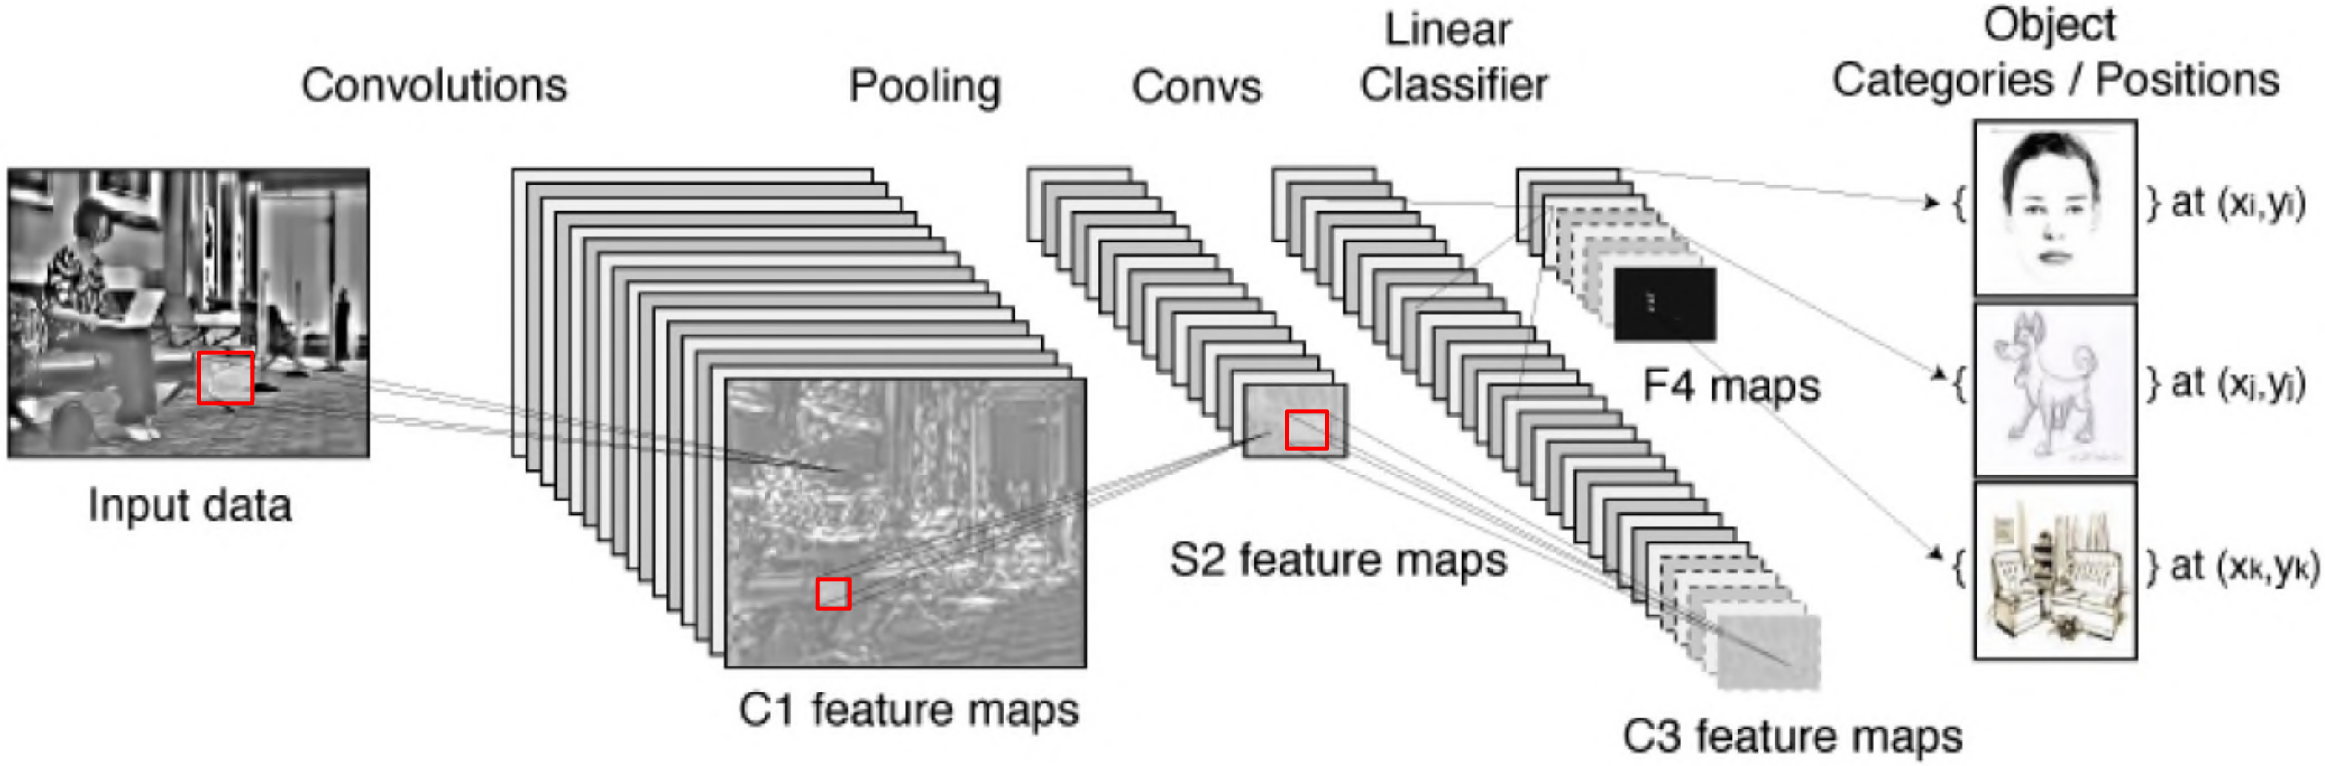

Slide courtesy: Nando de Freitas, University of Oxford

Convnets (Fukushima, LeCun, Hinton)

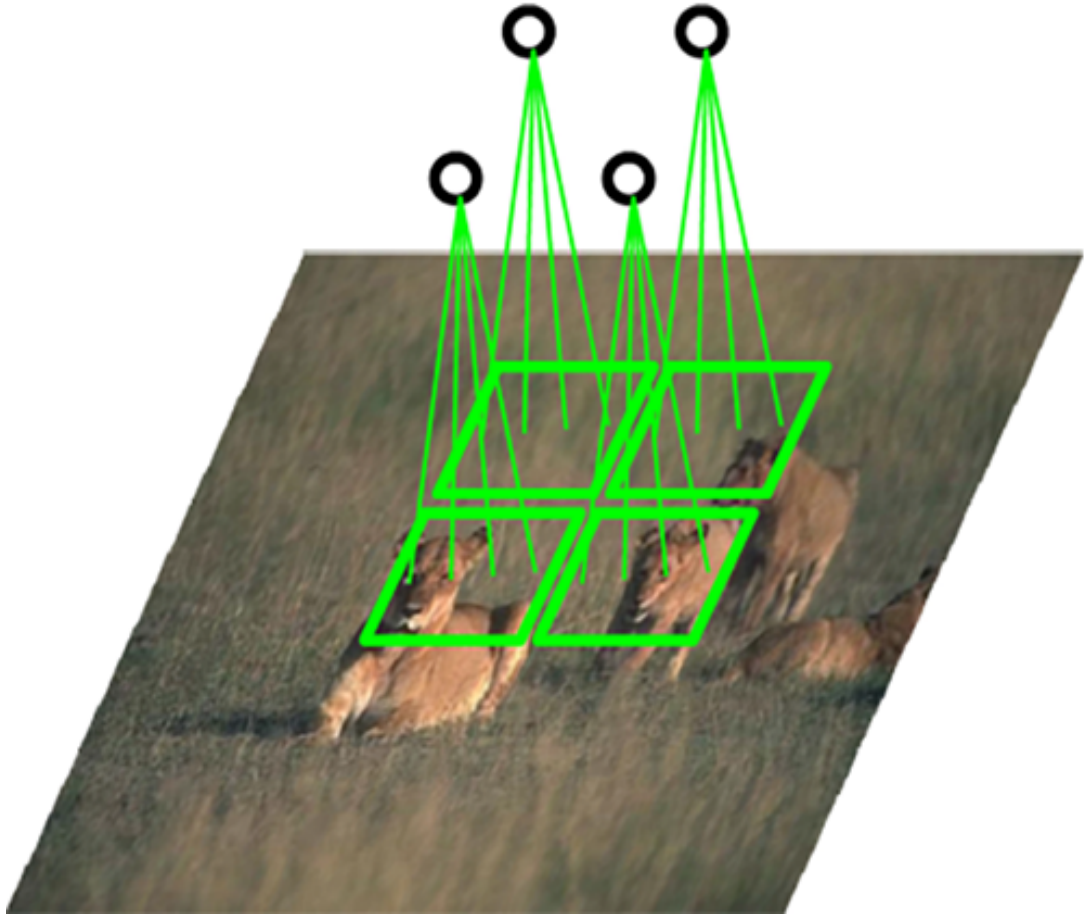

Slide courtesy: Nando de Freitas, University of Oxford

Pooling

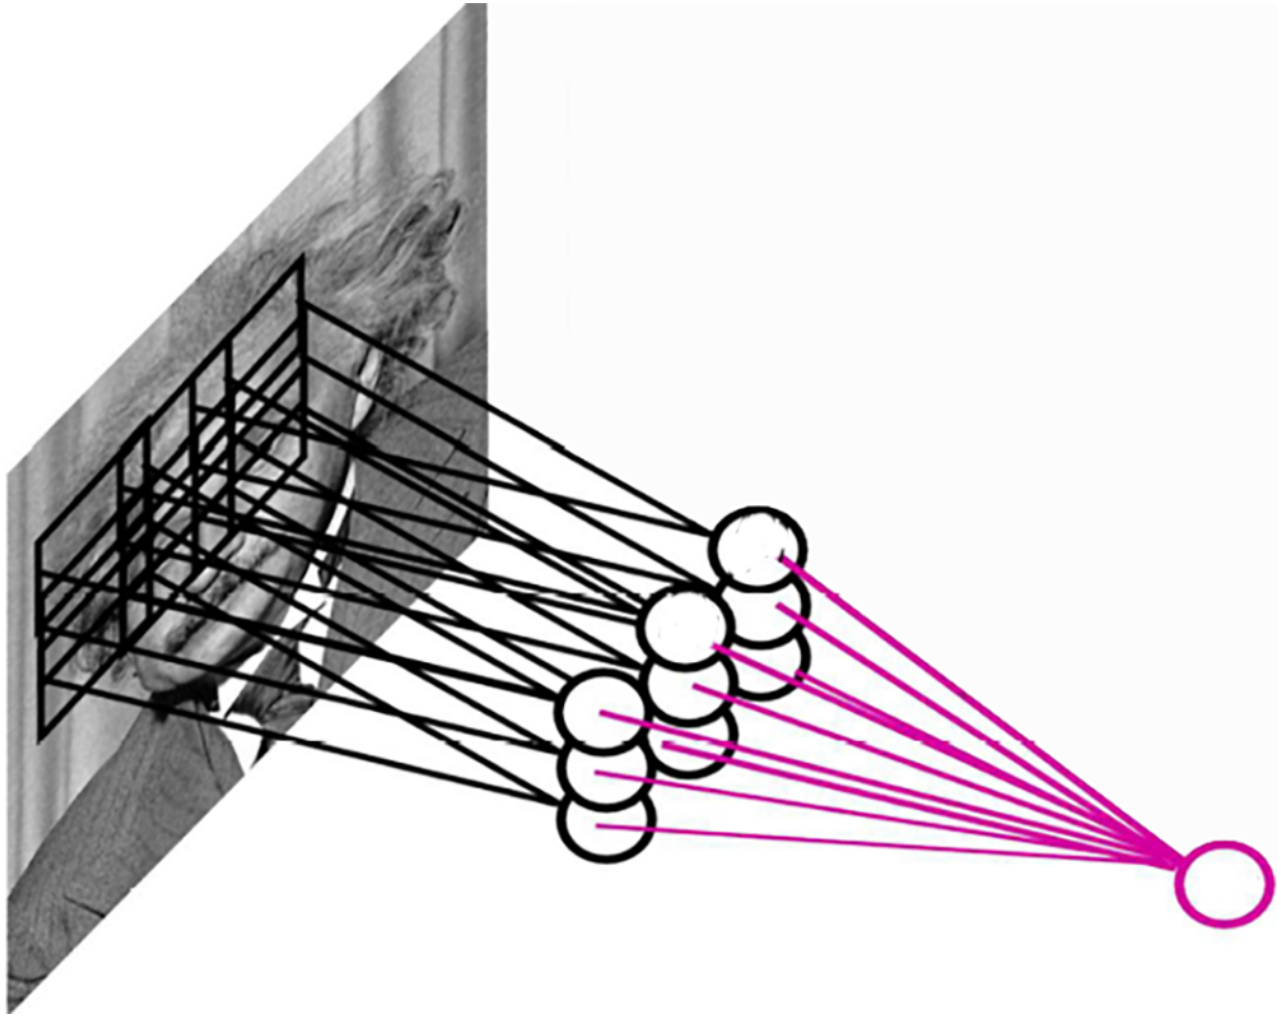

By "pooling" (e.g., max or average) filter responses, we gain robustness to the exact spatial location of features.

Image courtesy: Nando de Freitas, University of Oxford

Convolution

Filter

0	1	2
2	2	0
0	1	2

Image

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

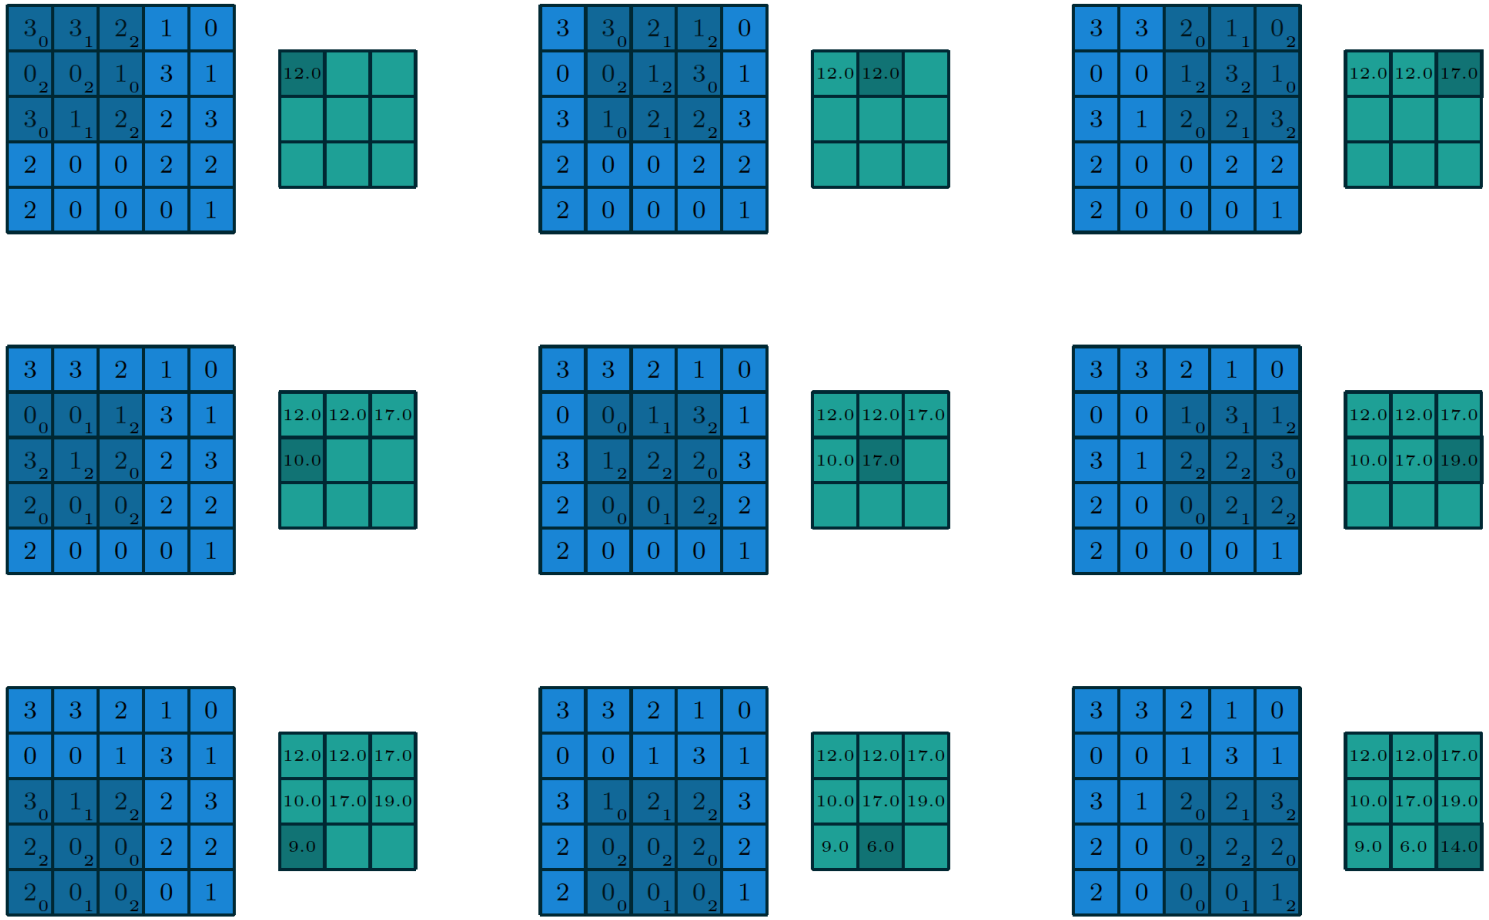

Image courtesy:Vincent Dumoulin

Convolution with Zero Padding

Filter

0	1	2
2	2	0
0	1	2

Image

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

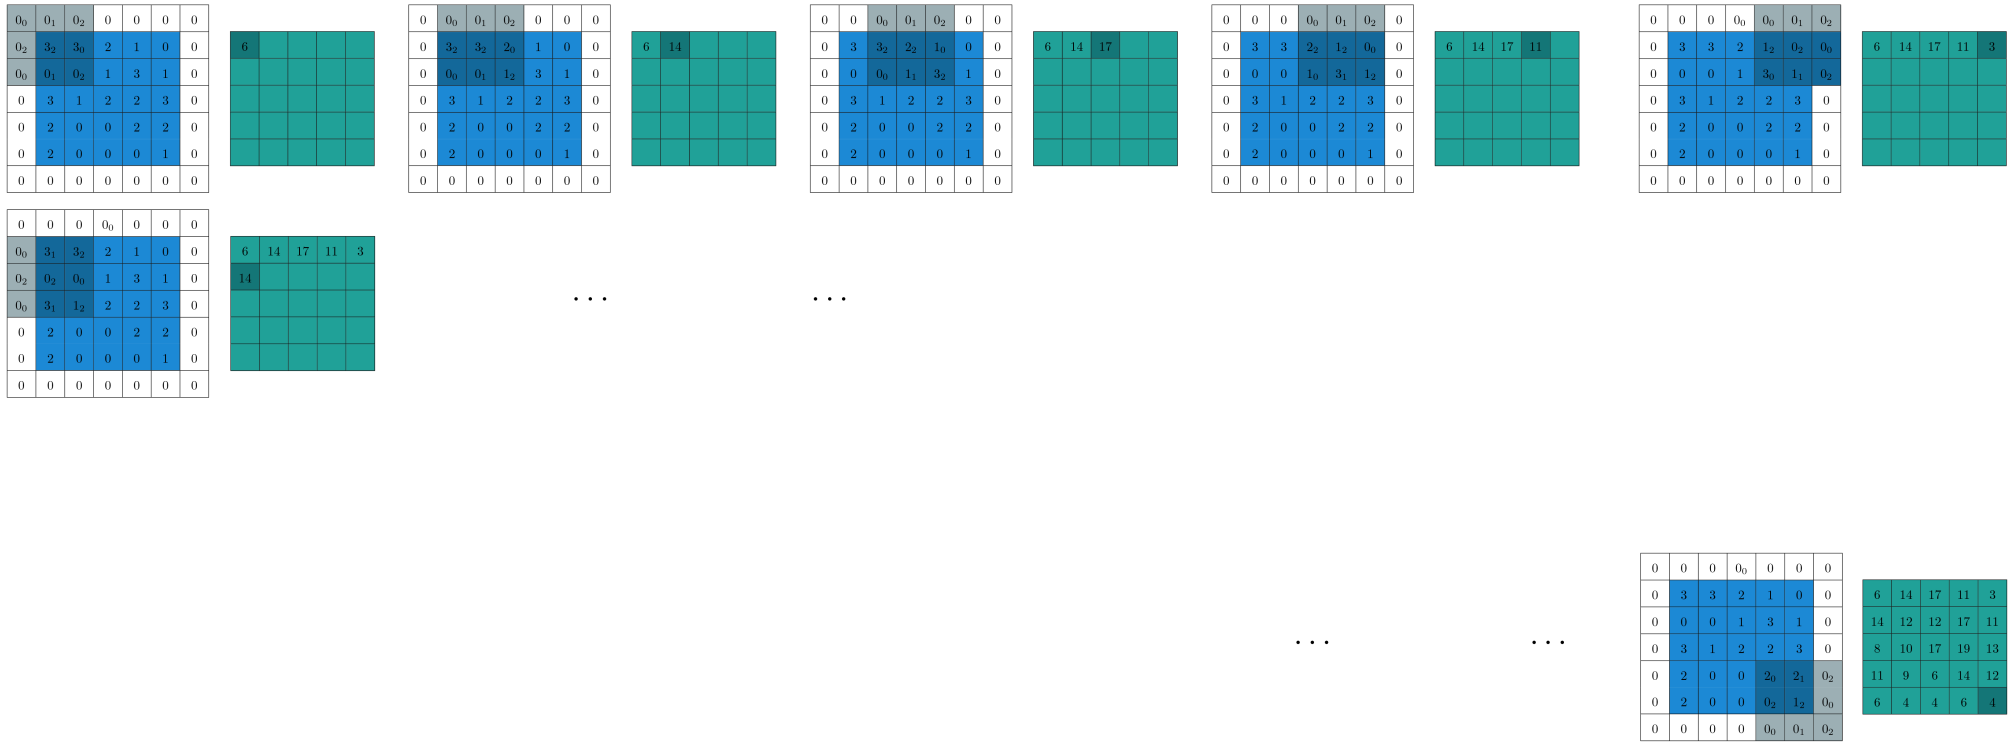

Convolution with Strides

Convolving a 3x3 kernel over a 5x5 input using 2x2 strides

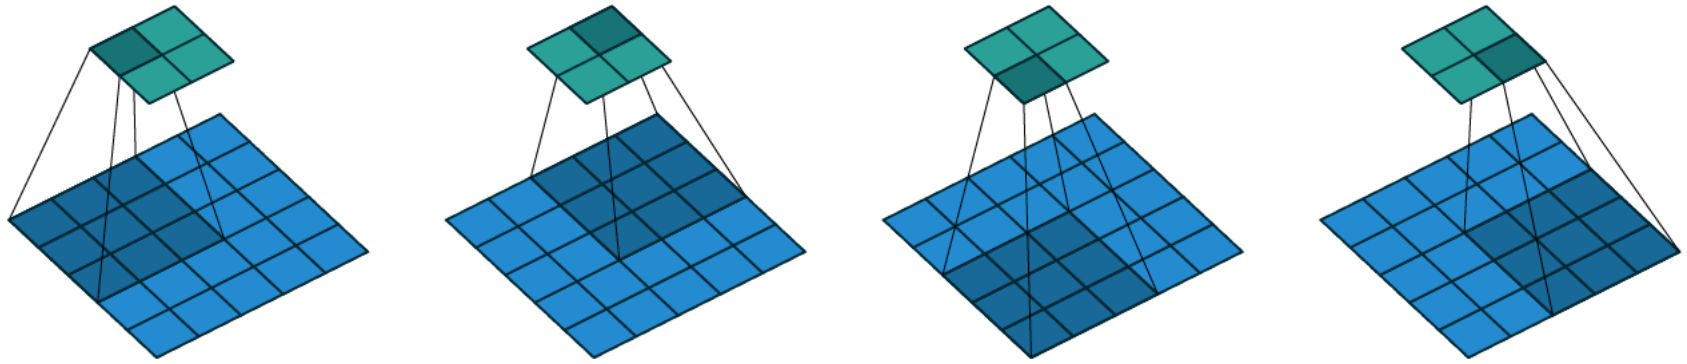

Image courtesy:Vincent Dumoulin

Convolution with Strides and Zero Padding

Convolving a 3x3 kernel over a 5 x5 input using 1 x1 strides and half padding

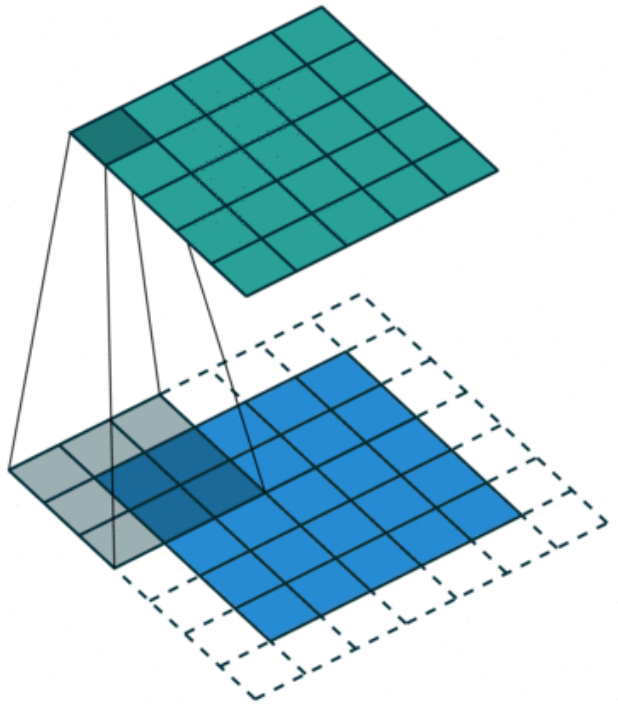

Image courtesy:Vincent Dumoulin

Inputs Generally have Multiple Channels

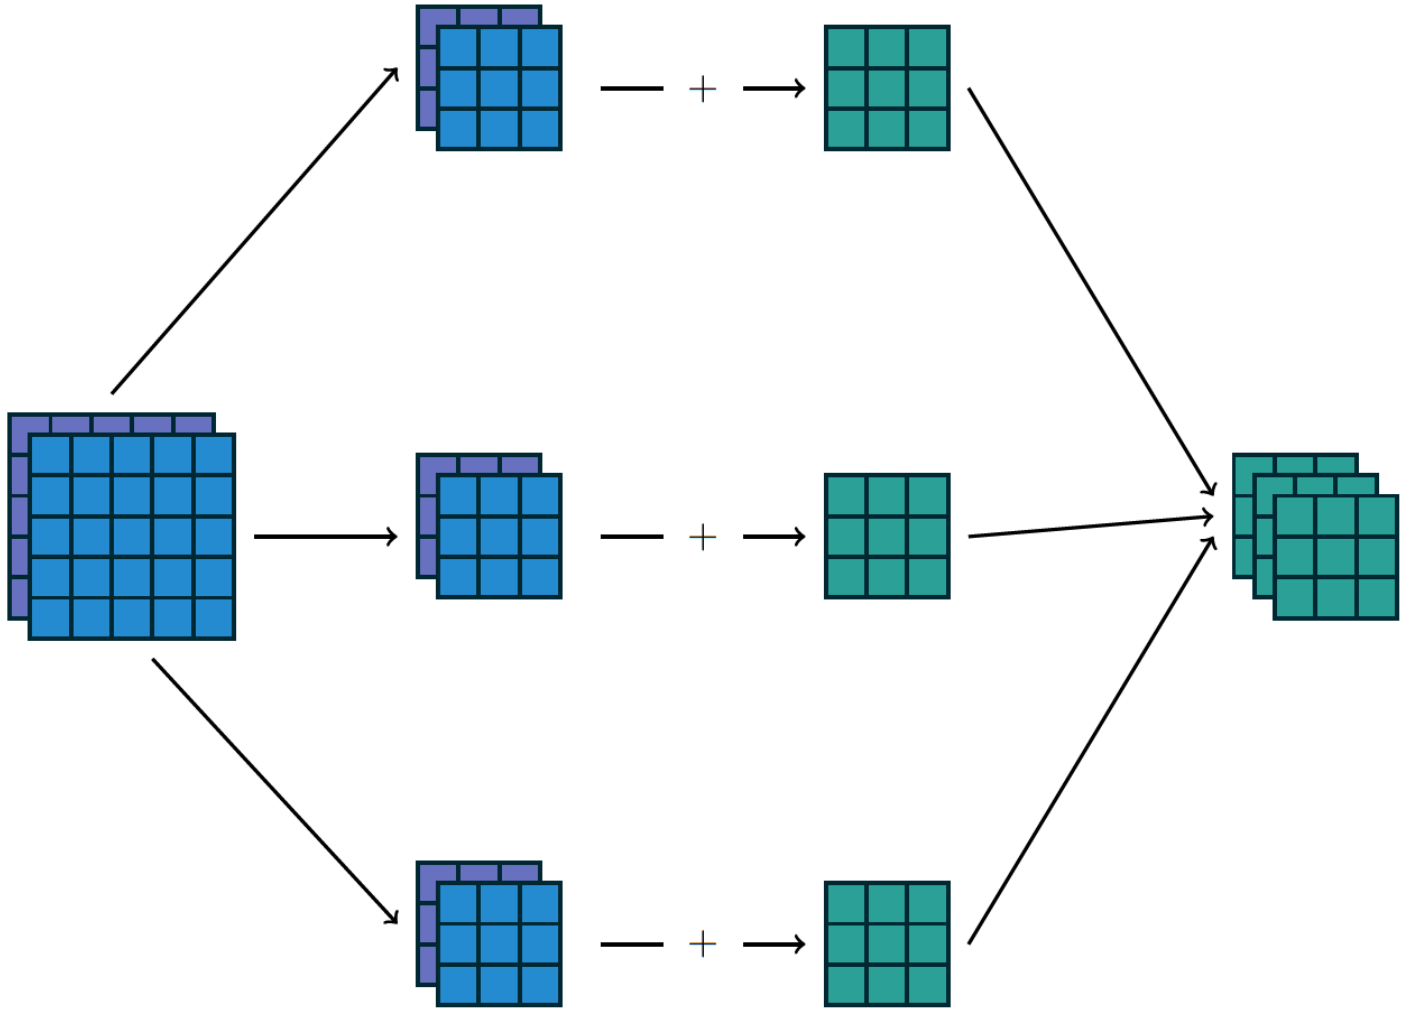

Image courtesy:Vincent Dumoulin

Convolution Arithmetic

(For simplicity we are assuming square image and filter/kernel)

Image width = image height = w

Filter width = Filter height = k

Stride = s

Output size = $\left\lfloor \frac{w-k}{s} \right\rfloor + 1$

Padding = p → This means image dimension becomes $w + 2p$

So, output size = $\left\lfloor \frac{w+2p-k}{s} \right\rfloor + 1$

Note the box function

$w = 6, p = 1, k = 3, s = 2$

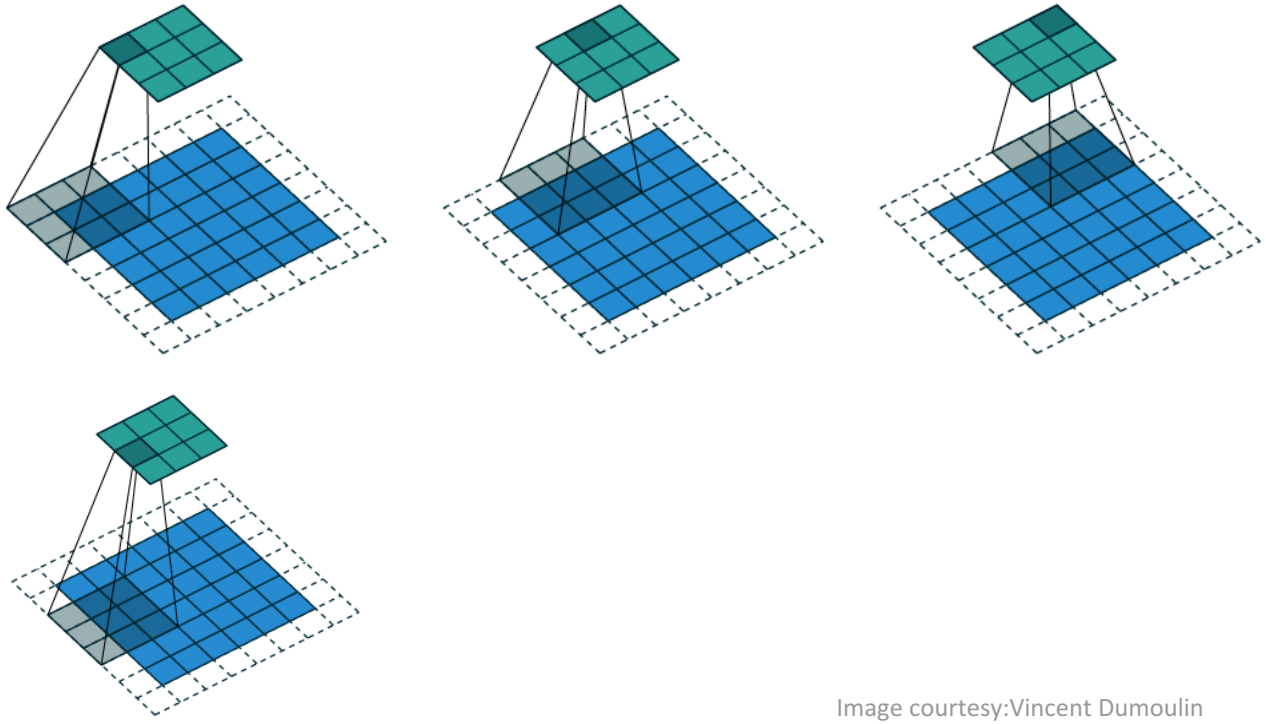

Image courtesy:Vincent Dumoulin

Convolution Arithmetic

Input volume: 32x32x3 [w, h, c].
64 filters of size 3x3 [k, k] with
stride 2 [s], pad 1 [p]

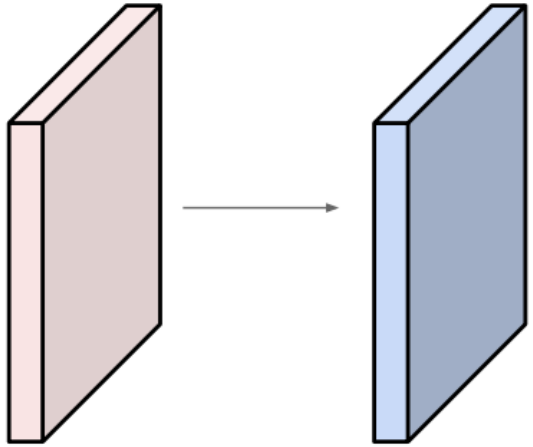

What is the output feature map size?

$$\left\lfloor \frac{32 + 2 * 1 - 3}{2} \right\rfloor + 1 = 16$$

So, $16 \times 16 \times 64$ [w, h, c]

And What is the number of parameters in this convolution layer?

$64 \times 3 \times 3 \times 3$ [c_out, w, h, c_in] = 1728

Image courtesy: Andrej Karpathy

Pooling

3x3 max-pooling on
5x5 input with
1x1 stride

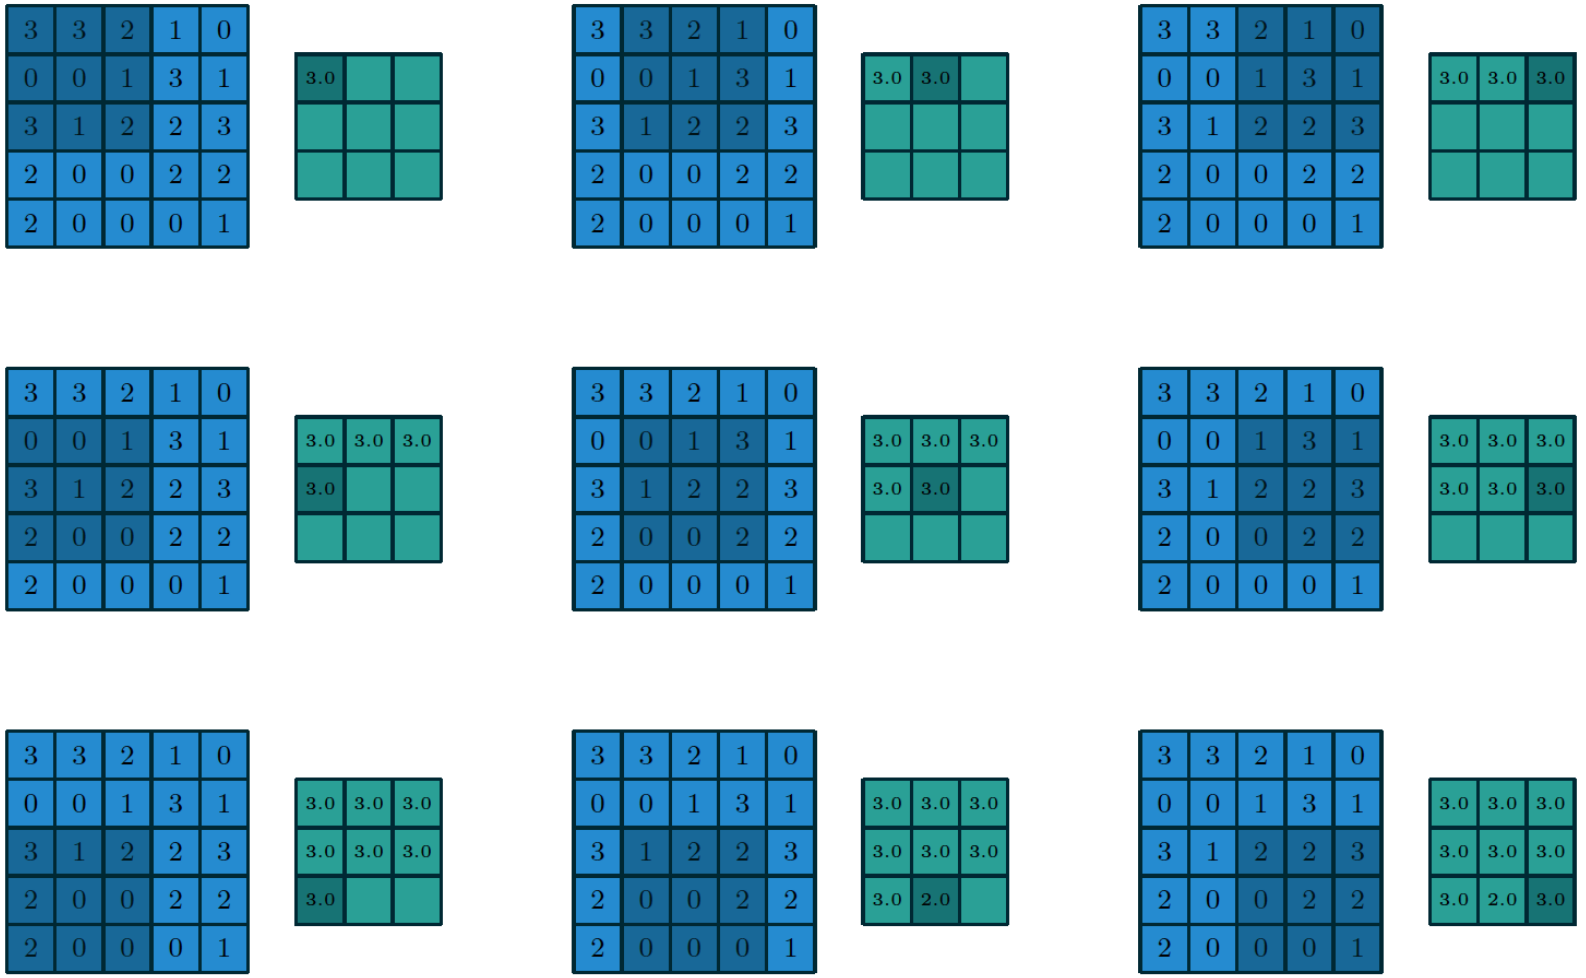

Image courtesy:Vincent Dumoulin

Pooling Arithmetic

(For simplicity we are assuming square input and max pooling kernel)

Input width = Input height = w

Filter width = Filter height = k

Stride = s

Output size = $\left\lfloor \frac{w-k}{s} \right\rfloor + 1$

Input volume: 32x32x3 [w, h, c].

Max-pooling kernel of size 2x2 [k, k] with stride 2 [s]

What is the output feature map size?

$\left\lfloor \frac{32-2}{2} \right\rfloor + 1 = 16$ So, 16×16×3 [w, h, c]

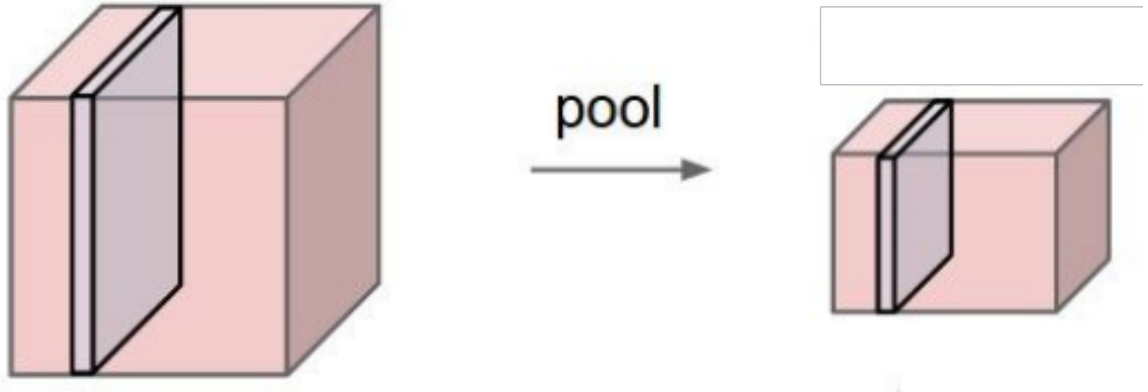

And What is the number of parameters in this pooling layer?

0

Image courtesy: Andrej Karpathy

Visualizations

AlexNet (2012)

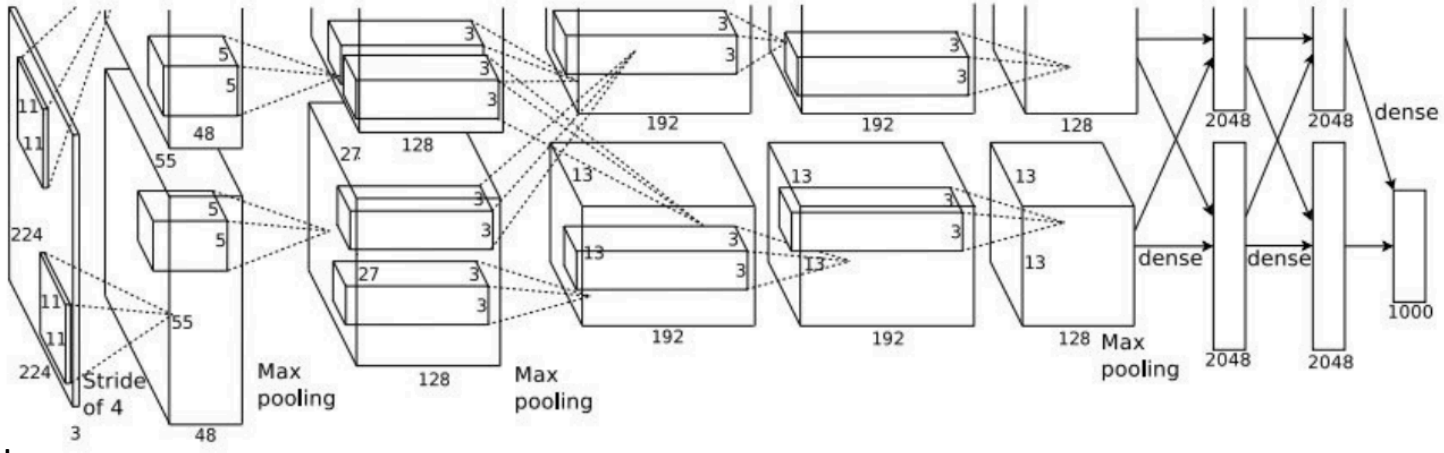

First layer (CONV1): 96 11x11 filters

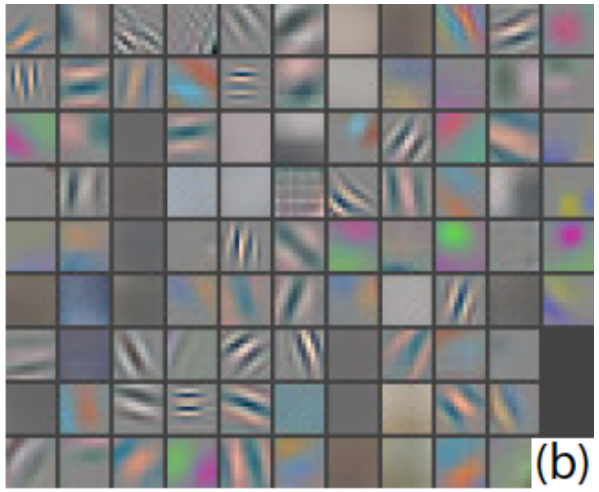

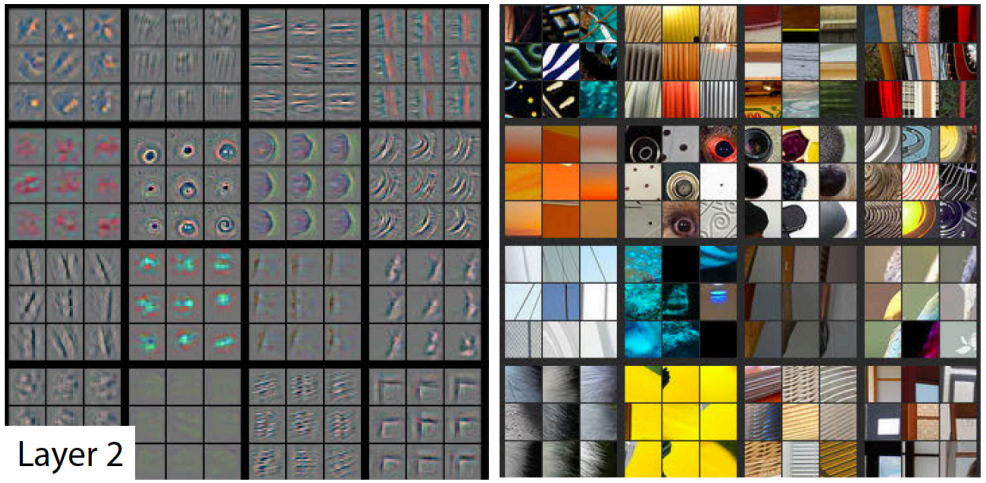

Layer 2

Image courtesy: Zeiler, Fergus, 2013

Visualizations

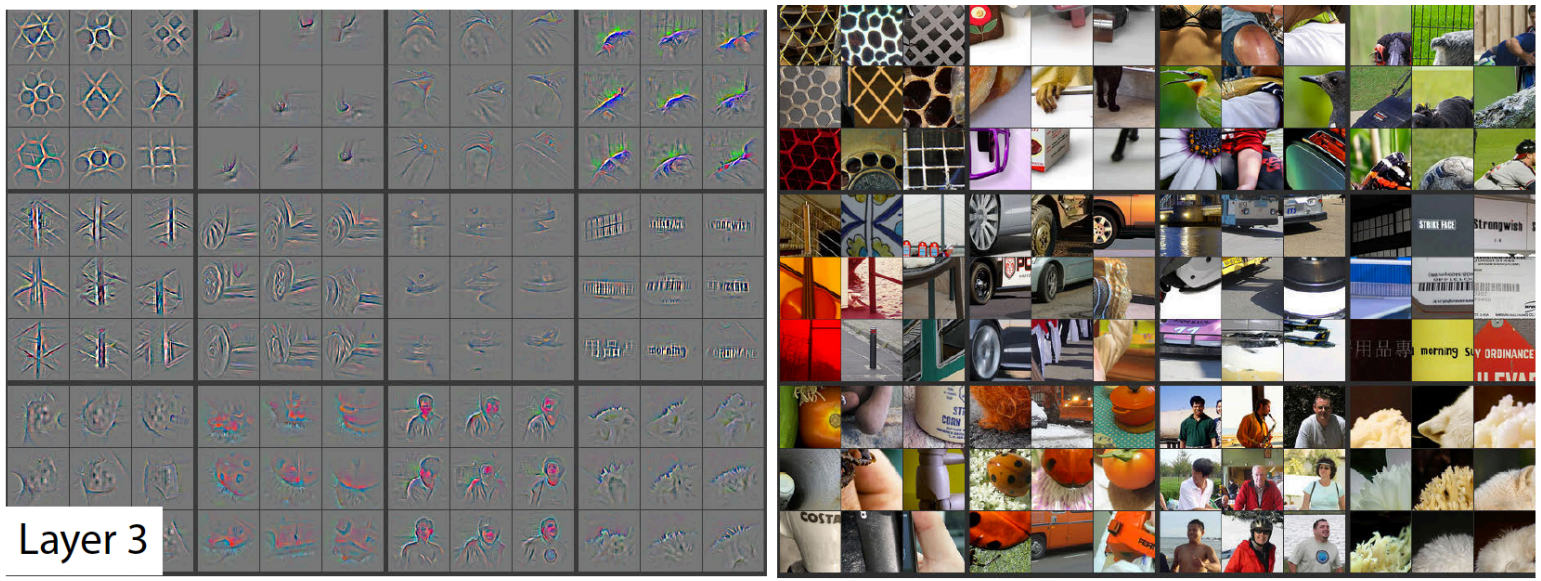

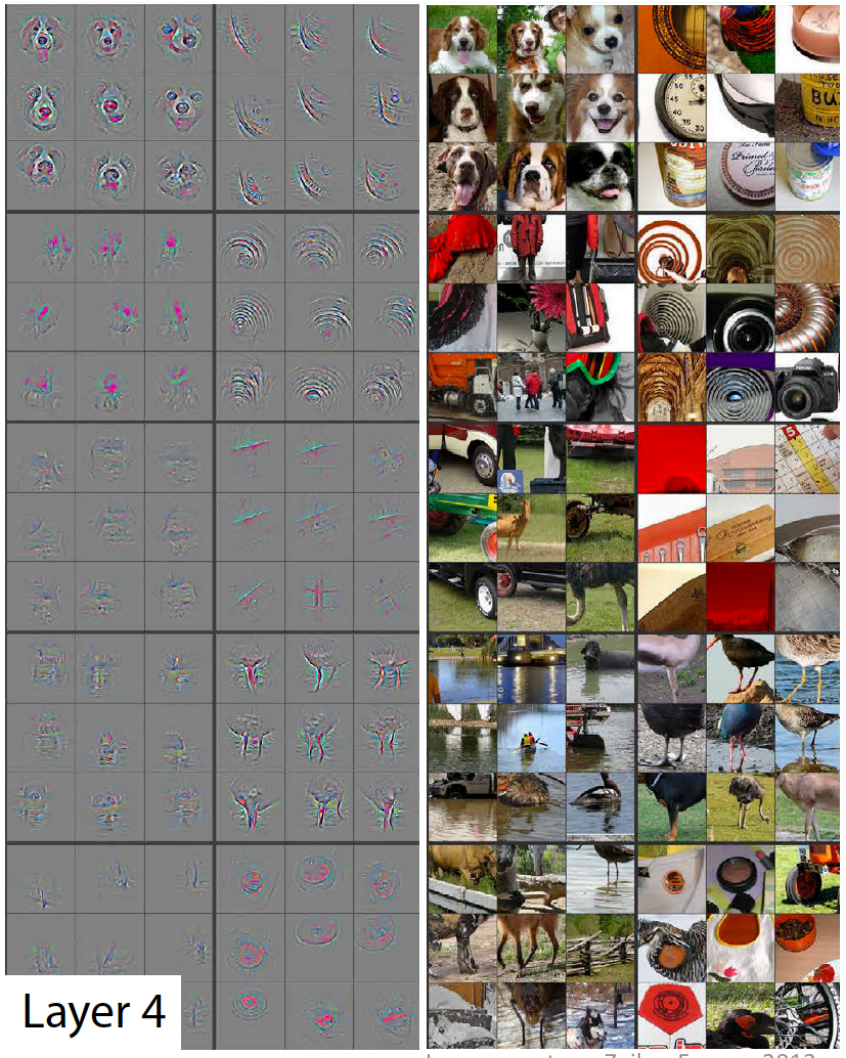

Image courtesy: Zeiler, Fergus, 2013

Distill ABOUT PRIZE SUBMIT

https://distill.pub/2018/building-blocks/

The Building Blocks of Interpretability

Interpretability techniques are normally studied in isolation.
We explore the powerful interfaces that arise when you combine them —
and the rich structure of this combinatorial space.